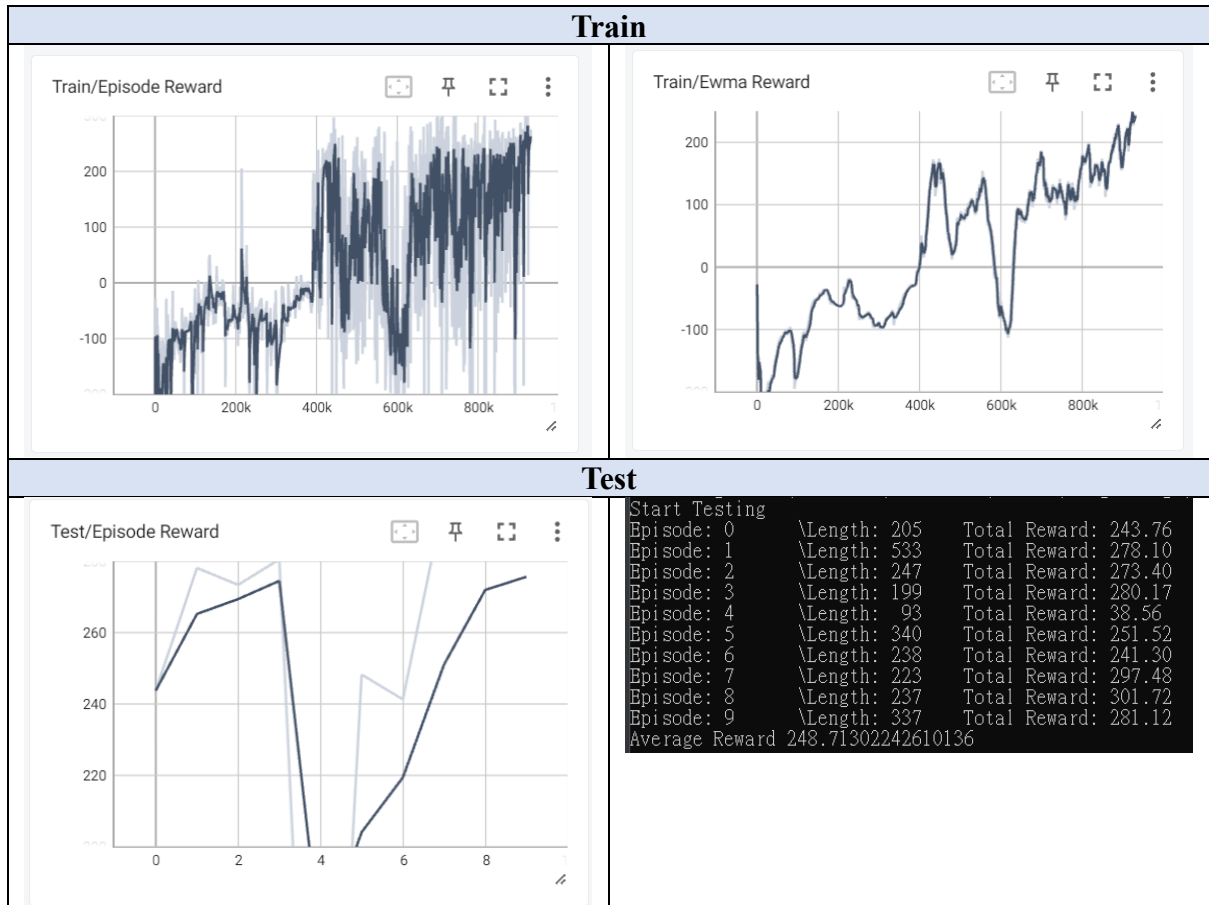


DL_LAB5_Deep Q-Network and Deep Deterministic Policy Gradient

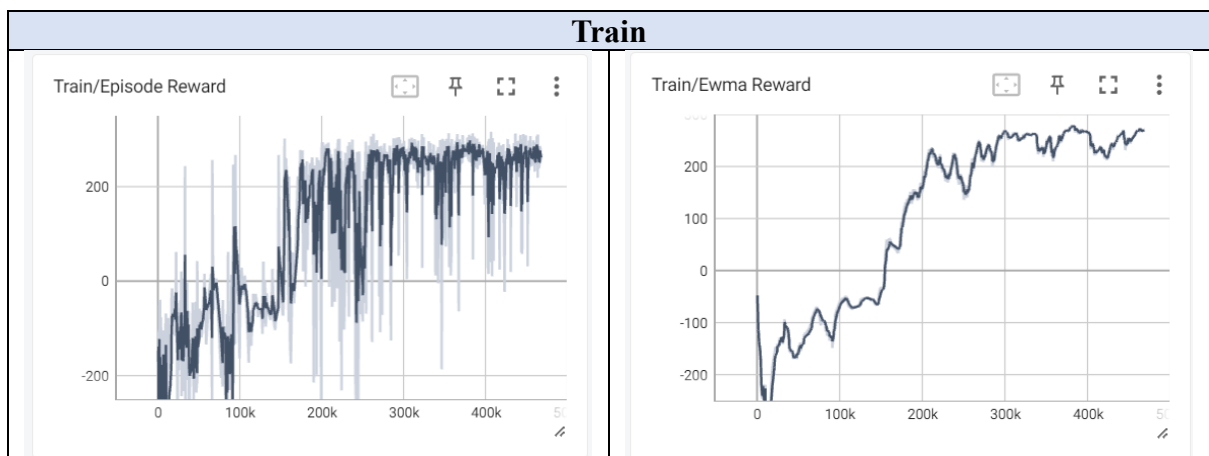
學號：312554004 姓名：林垣志

(1) Results of experiment:

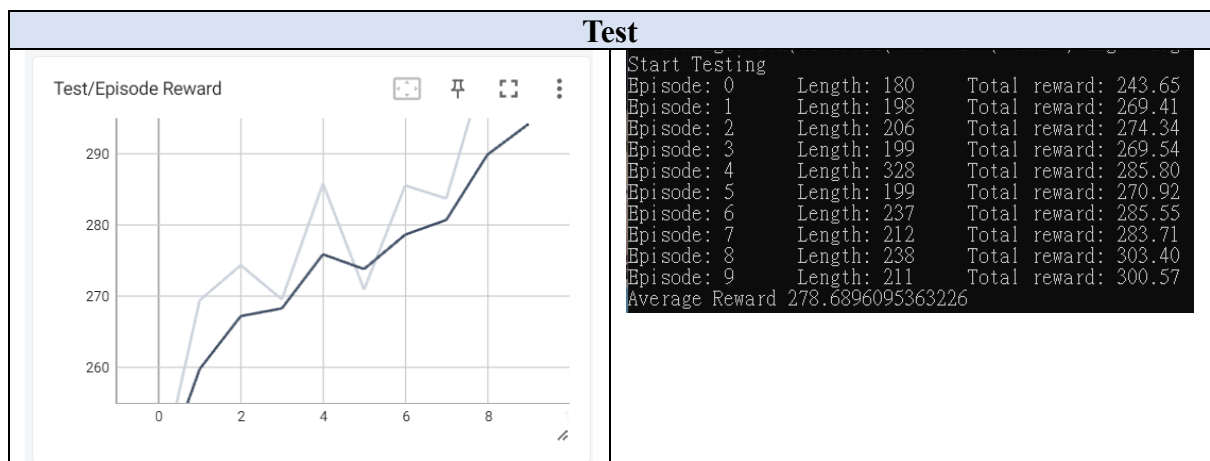
(a) LunarLander-v2 using DQN



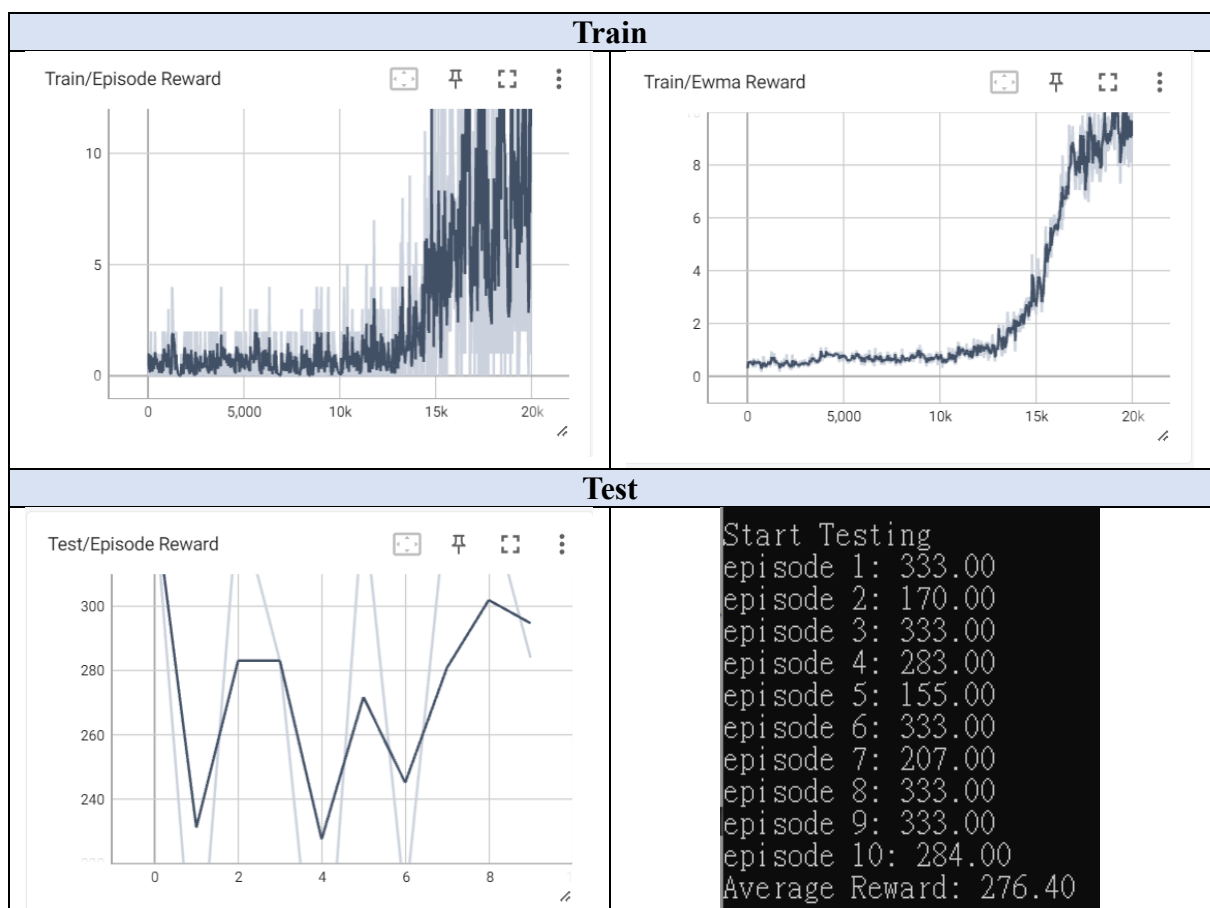
(b) LunarLanderContinuous-v2 using DDPG



(b) LunarLanderContinuous-v2 using DDPG

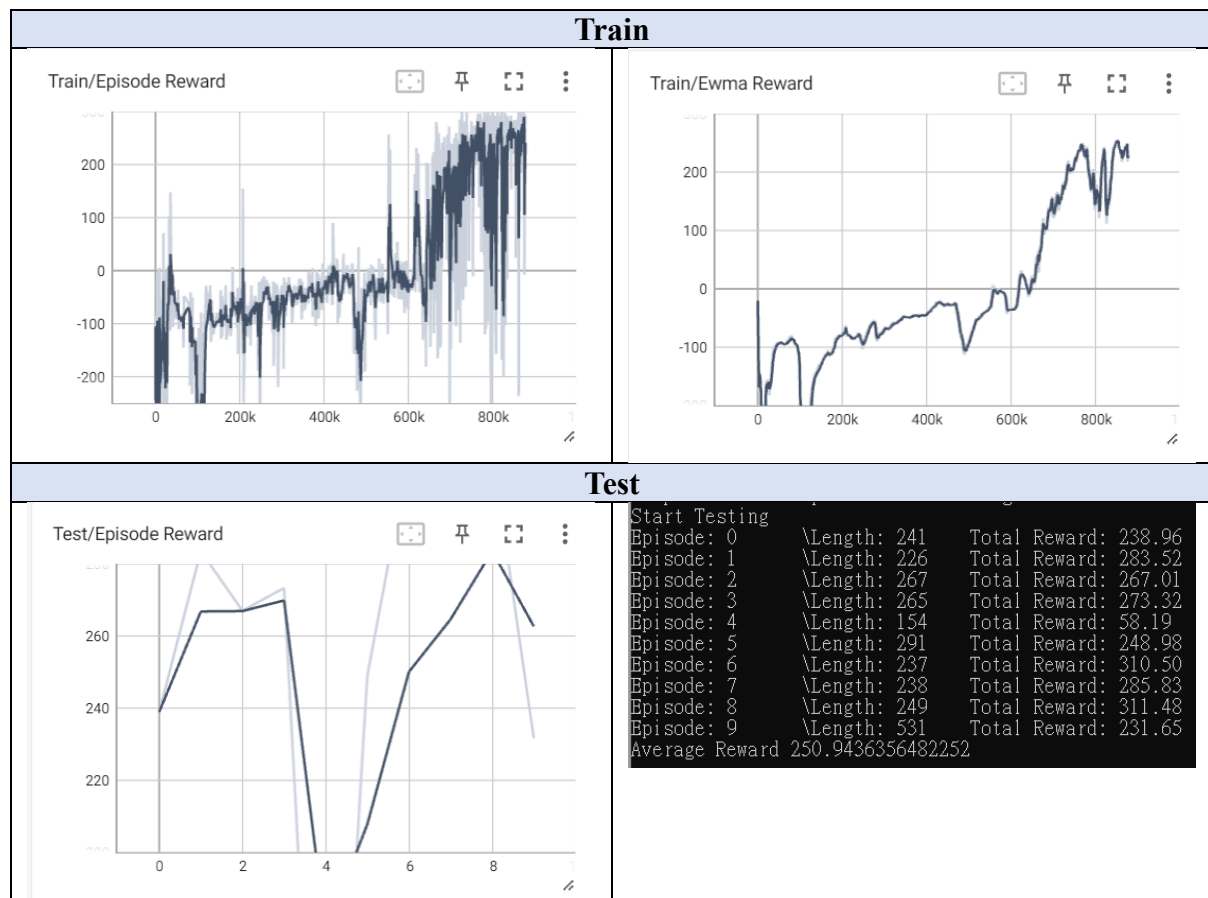


(c) BreakoutNoFrameskip-v4 using DQN



(2) Bonus:

LunarLander-v2 using DDQN



(3) Questions:

(a) Your implementation of Q network updating in DQN.

According to the algorithm, I will sample random minibatch of transition $(\phi_t, a_t, r_t, \phi_{t+1})$ from replay memory, then compute target value, and after that using MSE as loss function to update behavior net. In the `update_target_network` function, it's copy weights of behavior net to update target network every C steps.

```
def _update_behavior_network(self, gamma):
    # sample a minibatch of transitions
    state, action, reward, next_state, done = self._memory.sample(self.batch_size, self.device)

    """ TODO """
    # q_value = ?
    # with torch.no_grad():
    #     q_next = ?
    # q_target = ?
    # criterion = ?
    # loss = criterion(q_value, q_target)
    q_value = self._behavior_net(state).gather(1, action.long())
    with torch.no_grad():
        q_next = torch.max(self._target_net(next_state), 1)[0].view(-1, 1)
        q_target = reward + q_next * gamma * (1.0 - done)
    criterion = nn.MSELoss()
    loss = criterion(q_value, q_target)

    # optimize
    self._optimizer.zero_grad()
    loss.backward()
    nn.utils.clip_grad_norm(self._behavior_net.parameters(), 5)
    self._optimizer.step()

def _update_target_network(self):
    """ update target network by copying from behavior network """
    """ TODO """
    self._target_net.load_state_dict(self._behavior_net.state_dict())
```

(b) Your implementation and the gradient of actor updating in DDPG

According to the algorithm, it's same as DQN for the first, and let behavior actor network predict an action which can get maximum q value by critic network, using actor network to generate actions, and then get the q value by critic network, and after that computing the mean q value and back-propagation to update actor network. Therefore, we define actor loss as follows.

(c) Your implementation and the gradient of critic updating in DDPG

Like DQN, I will sample random minibatch from replay memory, and use actions of batch to calculate corresponding q value by critic network, also need to calculate target value. And after that, get action of next state by target actor network, generate the q value of next state by target critic network, multiply discount factor and plus reward, we can get the target q value. Finally, we can calculate MSE loss by q value and target q value.

```
def _update_behavior_network(self, gamma):
    actor_net, critic_net, target_actor_net, target_critic_net = self._actor_net,
                                                                self._critic_net, self._target_actor_net, self._target_critic_net

    actor_opt, critic_opt = self._actor_opt, self._critic_opt

    # sample a minibatch of transitions
    state, action, reward, next_state, done = self._memory.sample(self.batch_size, self.device)

    ## update critic ##
    # critic loss
    """ TODO """
    # q_value = ?
    # with torch.no_grad():
    #     a_next = ?
    #     q_next = ?
    #     q_target = ?
    # criterion = ?
    # critic_loss = criterion(q_value, q_target)
    q_value = critic_net(state, action)
    with torch.no_grad():
        q_next = target_critic_net(next_state, target_actor_net(next_state))
        q_target = reward + gamma * q_next * (1 - done)

    criterion = nn.MSELoss()
    critic_loss = criterion(q_value, q_target)

    # optimize critic
    actor_net.zero_grad()
    critic_net.zero_grad()
    critic_loss.backward()
    critic_opt.step()

    ## update actor ##
    # actor loss
    """ TODO """
    # action = ?
    # actor_loss = ?
    action = actor_net(state)
    actor_loss = -critic_net(state, action).mean()

    # optimize actor
    actor_net.zero_grad()
    critic_net.zero_grad()
    actor_loss.backward()
    actor_opt.step()
```